

## TTIC 31230 Fundamentals of Deep Learning

### Problem set 3

Due Thursday 11:59 pm, January 26

- Zip all your ipynb&pdf file with name PS3\_yourfullname to: ttic.dl.win.2017@gmail.com.
- Late Submission: submitting late work will be penalized 10% per day, maximum three days delay allowed, no submission allowed after that.

This problem sets involves writing convolution and maxpool layers in EDF. There are no analytical problems — just the coding. Your solutions should be entered into the appropriate cells of the jupyter notebook provided and you should receive an accuracy of around 40% on CIFAR 10.

**Part 1.** We want you to write a class MaxPool such that for

$$y = \text{MaxPool}(x, s)$$

the instance  $y$  is specified as follows.

Here  $x.value$  is the input image and is assumed to have shape  $(B, H, W, C)$  where  $B$  is the batch size,  $H$  is the image height,  $W$  is the image width, and  $C$  is the number of channels at each pixel position. The stride  $s$  is assumed to be an integer.

The shape of  $y.value$  should be

$$(B, \lfloor H/s \rfloor, \lfloor W/s \rfloor, C).$$

We want  $y.forward()$  to set  $y.value$  to the following tensor where  $u$  and  $v$  range over 0 to  $s - 1$ .

$$y.value[b, i, j, c] = \max_{u, v} x.value[b, si + u, sj + v, c]$$

If the maximum is tied at several spatial locations you can assume that the gradient with respect to each such maximal value equals the corresponding pooled value of  $y.grad$ . This assumption simplifies the backward method.

**Part 2.** This is the same as part 1 except that you are to build AvePool such that for

$$y = \text{AvePool}(x, s)$$

the instance  $y$  is specified as follows.

$$y.value[b, i, j, c] = \frac{1}{s^2} \sum_{u, v} x.value[b, si + u, sj + v, c]$$

**Part 3.** We want you to write a class Conv such that for

$$y = \text{Conv}(x, f, s, p)$$

the instance  $y$  is specified as follows.

Here  $x.value$  is the input image and is assumed to have shape  $(B, H, W, C)$  where  $B$  is the batch size,  $H$  is the image height,  $W$  is the image width, and  $C$  is the number of channels at each pixel position. Here  $f$  is assumed to be a square filter parameter with shape  $(K, K, C, C')$  where  $K$  is the spatial dimension of the filter. There is no batch dimension for parameters. The stride  $s$  and the padding  $p$  are assumed to be integers.

In the forward method one must first pad the  $x.value$  with padding width  $p$ . Let  $x'$  be the result of padding.  $x'$  should have shape  $(B, H + 2p, W + 2p, C)$  where we have

$$x'[:, p : p + H, p : p + W, :] = x.value$$

and all other values of  $x'$  are zero.

For a square filter  $K \times K$  filter the shape of  $y.value$  should be

$$(B, \lfloor (H + 2p - K)/s \rfloor + 1, \lfloor (W + 2p - K)/s \rfloor + 1, C').$$

We want  $y.forward()$  to set  $y.value$  to the following tensor where  $u$  and  $v$  range over 0 to  $K - 1$  (the spatial coordinates of the filter).

$$y.value[b, i, j, c'] = \left( \sum_{u, v, c} x'[b, si + u, sj + v, c] f.value[u, v, c, c'] \right)$$

Hint: you can write a Python loop over  $i$  and  $j$  (but not  $b$ ) and for  $i, j$  fixed use `np.matmul` to do the summation over  $u, v$  and  $c$ .

In this problem set each batch component of  $\ell.grad$  is initialized to  $1/B$  where  $B$  is the batch size rather than being initialized to 1. Each backward method adding to the gradient of a parameter then sums over the batch rather than average.

**Part 4.** Once you have constructed the Conv and MaxPool layers construct and test the following model.

1. A Conv Layer with a  $3 \times 3$  filter mapping 3 color channels to 32 feature channels and with stride 1, padding 1, and a ReLU activation function resulting in a  $32 \times 32$  image.
2. A MaxPool layer with stride 4 resulting in image dimensions  $8 \times 8$  with 32 channels.
3. A Conv layer with a  $3 \times 3$  filter mapping 32 channels to 64 channels and with padding 0 and stride 1 and a ReLU activation function resulting in a  $6 \times 6$  image with 64 channels.
4. An AvePool layer with stride 6 resulting in a  $1 \times 1$  image with 64 channels.

5. A Conv layer on the  $1 \times 1$  image with a  $1 \times 1$  filter mapping 64 channels to 10 channels with a ReLU nonlinearity. This is functionally equivalent to a perceptron layer but you can use your implementation of the Conv layer.
6. Reshape the  $1 \times 1$  image with 10 channels to a 10 dimensional vector using the Reshape class in edf.py.
7. Softmax, LogLoss etc, which are standard in classification task.

Filter parameters should be constructed using `edf.Param(edf.xavier(shape))` which provides a well-initialized random value for a filter tensor with the given shape. After Convolution, you should add a bias vector  $\beta$  using `edf.Add` before each ReLU activation. Bias vector parameters should be initialized to zero. For

$$z = \text{Add}(y, \beta)$$

with  $y$  having shape  $(B, H, W, C)$  we have

$$z.\text{value}[b, i, j, c] = y.\text{value}[b, i, j, c] + \beta[c]$$

The notebook provides softmax and logloss layers to define the loss and code to run the model.

Partial credit will be give for code that is “almost right”.